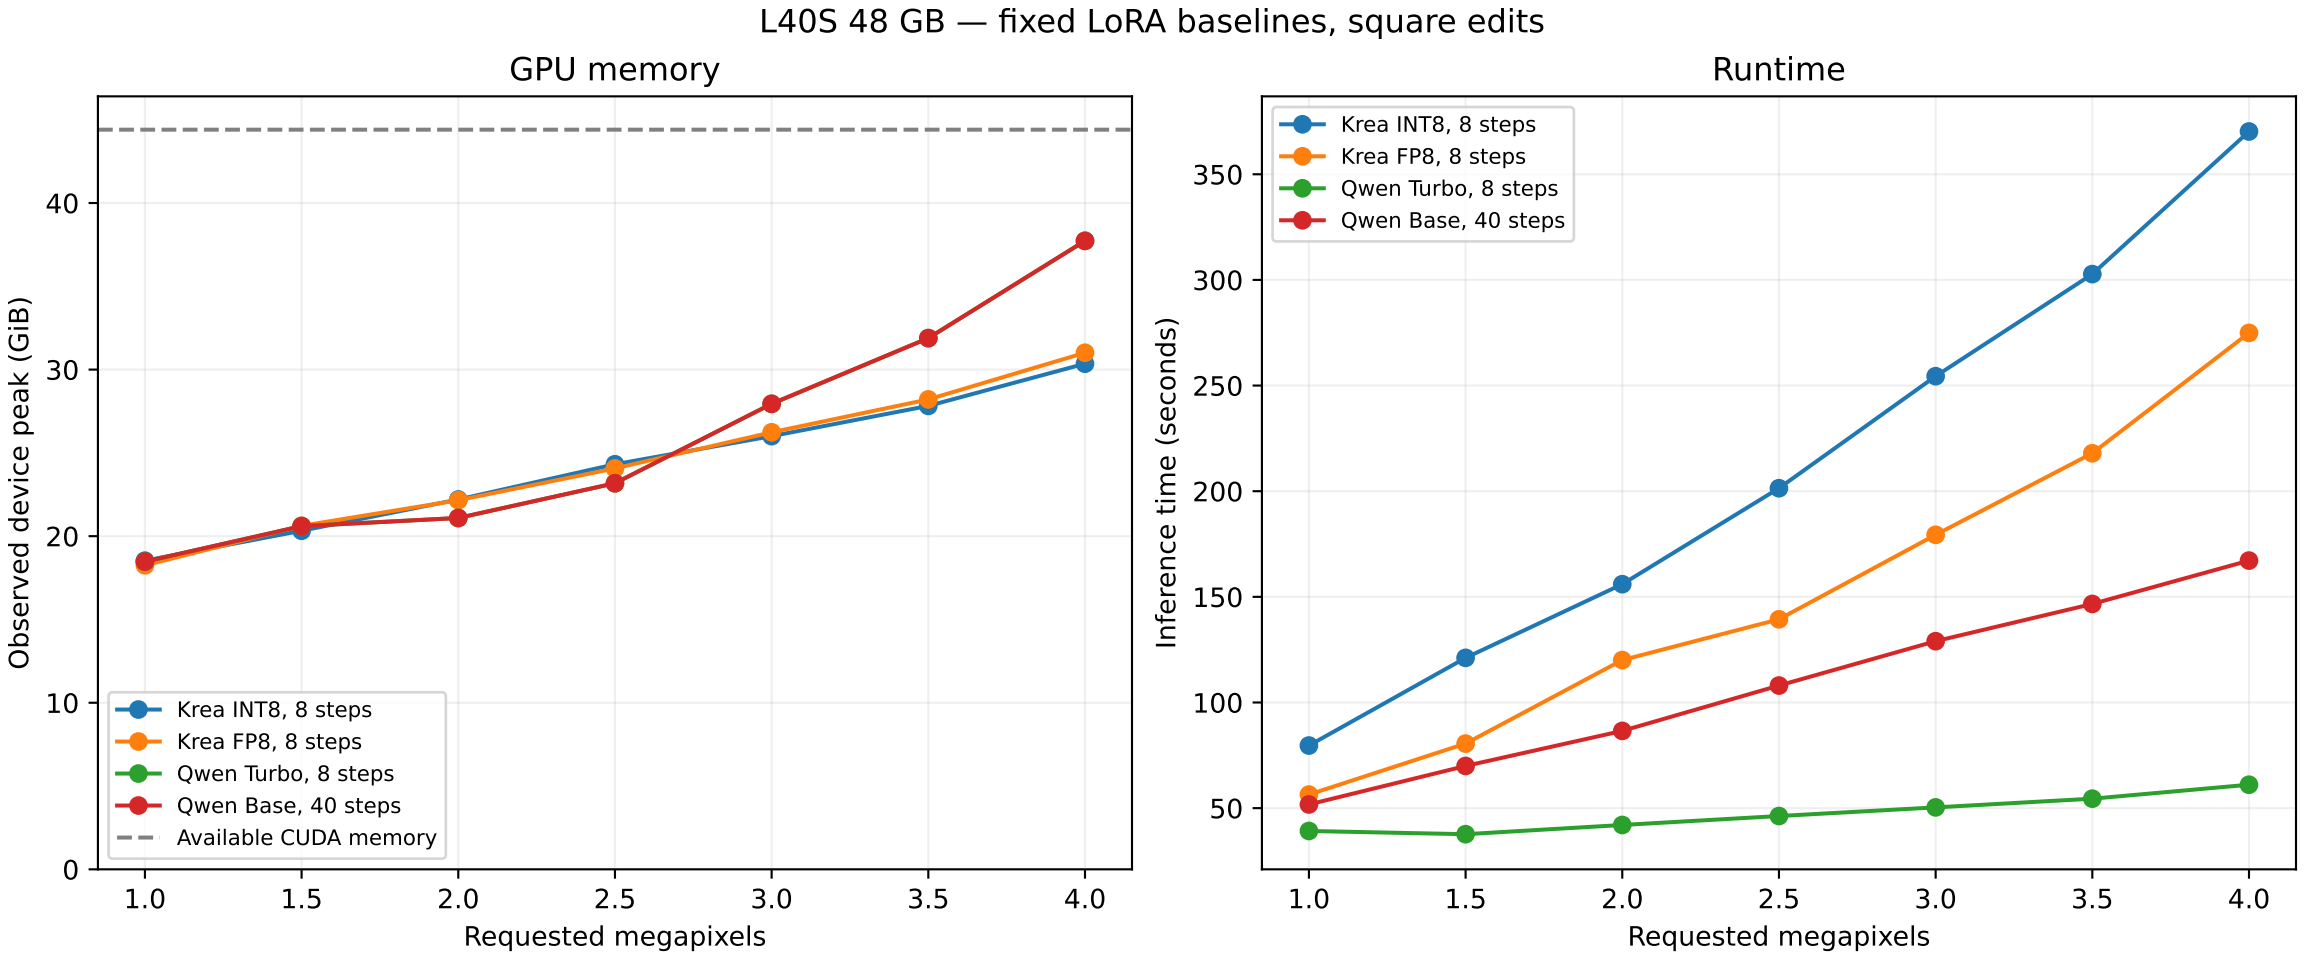
L40S 48 GB — fixed LoRA baselines, square edits
GPU memory
Observed device peak (GiB)
Requested megapixels
Krea INT8, 8 steps
Krea FP8, 8 steps
Qwen Turbo, 8 steps
Qwen Base, 40 steps
Available CUDA memory
Runtime
Inference time (seconds)
Requested megapixels
Krea INT8, 8 steps
Krea FP8, 8 steps
Qwen Turbo, 8 steps
Qwen Base, 40 steps