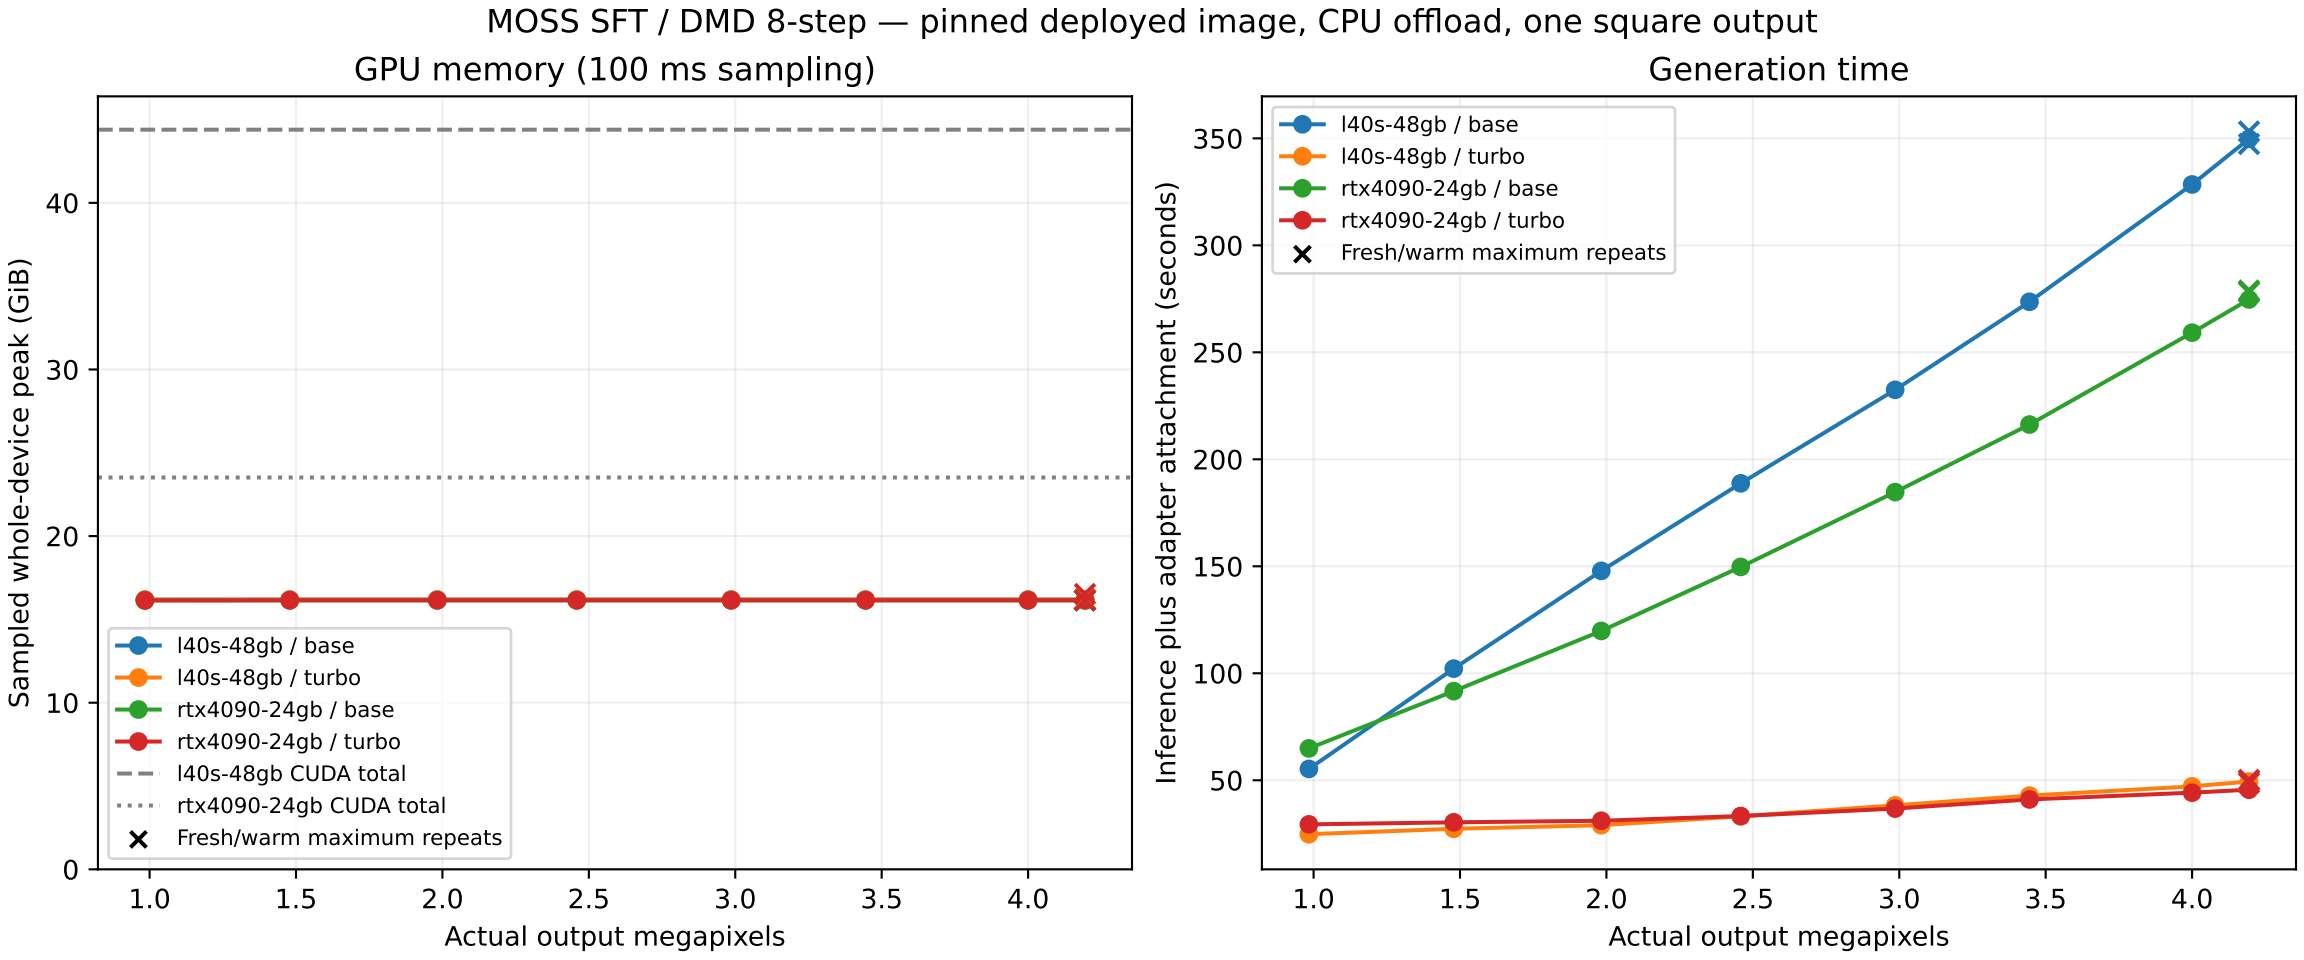
MOSS SFT / DMD 8-step — pinned deployed image, CPU offload, one square output
GPU memory (100 ms sampling)
Sampled whole-device peak (GiB)
Actual output megapixels
l40s-48gb / base
l40s-48gb / turbo
rtx4090-24gb / base
rtx4090-24gb / turbo
l40s-48gb CUDA total
rtx4090-24gb CUDA total
Fresh/warm maximum repeats
Generation time
Inference plus adapter attachment (seconds)
Actual output megapixels
l40s-48gb / base
l40s-48gb / turbo
rtx4090-24gb / base
rtx4090-24gb / turbo
Fresh/warm maximum repeats